

CS 5/7320

Artificial Intelligence

Logical Agents

AIMA Chapter 7

Slides by Michael Hahsler
based on slides by Svetlana Lazepnik
with figures from the AIMA textbook



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).

["Exercise Plays Vital Role Maintaining Brain Health"](#)
by [A Health Blog](#)

What is logic?



Logic is a formal system for manipulating facts so that true conclusions may be drawn



Syntax: rules for constructing valid sentences

E.g., $x + 2 \geq y$ is a valid arithmetic sentence, $\geq x^2 y +$ is not

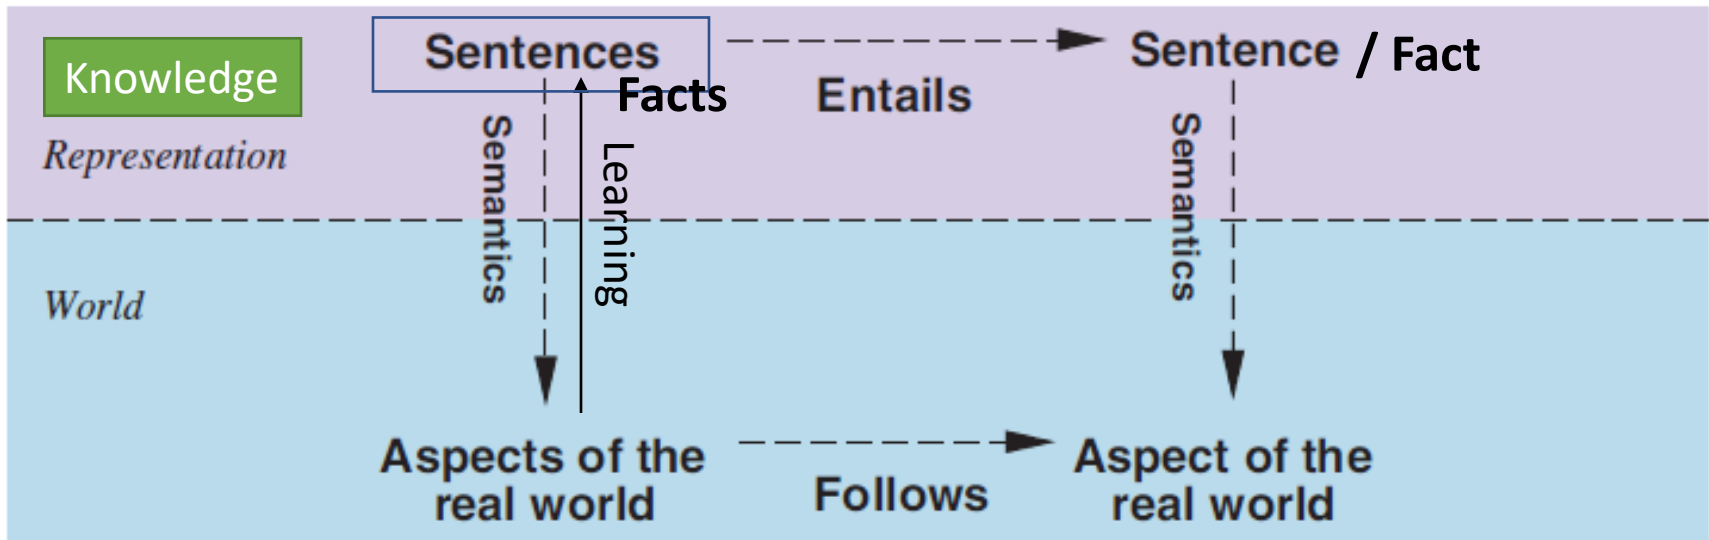


Semantics: “meaning” of sentences, or relationship between logical sentences and the real world

Specifically, semantics defines truth of sentences

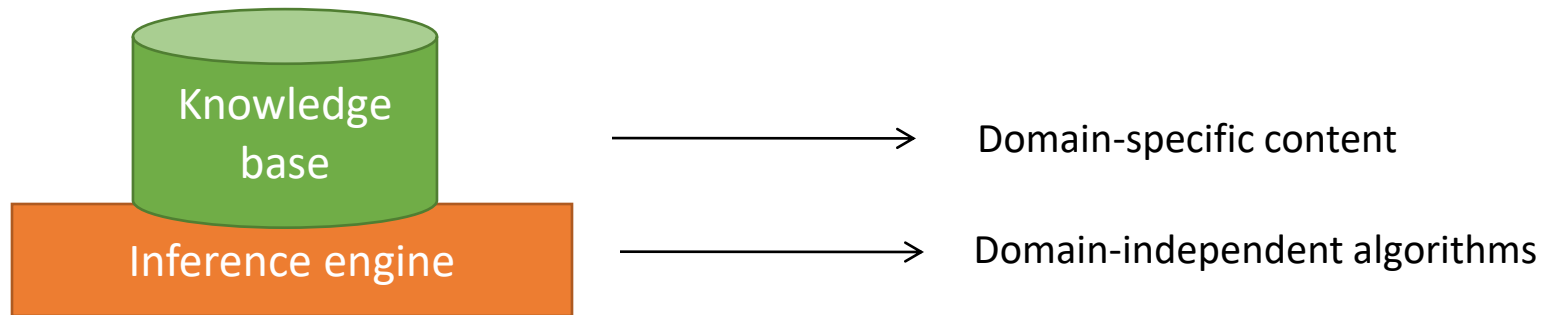
E.g., $x + 2 \geq y$ is true in a world where $x = 5$ and $y = 7$

Relation between Knowledge Representation and the World



- **Facts:** Sentences we know to be true.
- **Possible worlds:** all worlds/models which are consistent with the facts we know (compare with belief state).
- Learning new facts reduces the number of possible worlds.
- **Entailment:** A sentence logically follows from what we already know.

Knowledge-based agents



- Knowledge base (KB) = set of **sentences** in a **formal** language (knowledge representation) that are known to be true = **set of facts**
- **Declarative** approach to building an agent: Define what it needs to know
- Distinction between data (knowledge) and program (inference)
- Fullest realization of this philosophy was in the field of expert systems or knowledge-based systems in the 1970s and 1980s

Generic Knowledge-based Agent

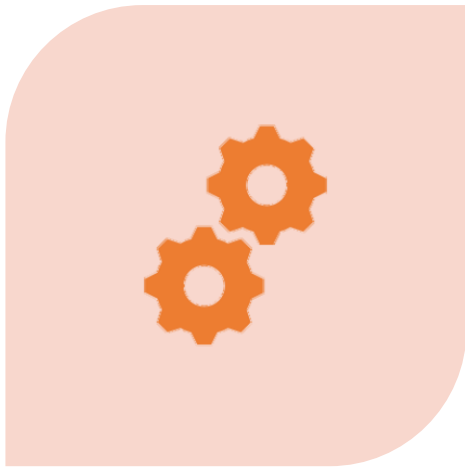
```
function KB-AGENT(percept) returns an action  
  persistent: KB, a knowledge base  
               t, a counter, initially 0, indicating time  
  
  TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))  
  action  $\leftarrow$  ASK(KB, MAKE-ACTION-QUERY(t))  
  TELL(KB, MAKE-ACTION-SENTENCE(action, t))  
  t  $\leftarrow$  t + 1  
  return action
```

Memorize percept
at time *t*

Ask for logical
action

Record action
taken at time *t*

Overview



PROPOSITIONAL LOGIC



FIRST ORDER LOGIC



Propositional Logic

Propositional logic: Syntax in BN-Form

Sentence \rightarrow *AtomicSentence* | *ComplexSentence*

AtomicSentence \rightarrow *True* | *False* | *P* | *Q* | *R* | ... = Symbols

ComplexSentence \rightarrow (*Sentence*)

| \neg *Sentence*

Negation

| *Sentence* \wedge *Sentence*

Conjunction

| *Sentence* \vee *Sentence*

Disjunction

| *Sentence* \Rightarrow *Sentence*

Implication

| *Sentence* \Leftrightarrow *Sentence*

Biconditional

OPERATOR PRECEDENCE : $\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow$

Validity and Satisfiability

A sentence is **valid** if it is true in **all** models (called a tautology)

e.g., *True*, $A \vee \neg A$, $A \Rightarrow A$, $(A \wedge (A \Rightarrow B)) \Rightarrow B$
useful to deduct new sentences.

A sentence is **satisfiable** if it is true in **some** model

e.g., $A \vee B$, C
useful to find new facts that satisfy all possible worlds.

A sentence is **unsatisfiable** if it is true in no models

e.g., $A \wedge \neg A$

Possible Worlds, Models and Truth Tables

A **model** specifies a “possible world” with the true/false status of each proposition symbol in the knowledge base

- E.g., **P** is true and **Q** is true
- With two symbols, there are $2^2 = 4$ possible worlds/models, and they can be enumerated exhaustively using:

A **truth table** specifies the truth value of a composite sentence for each possible assignment of truth values to its atoms. Each row is a model.

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	$P \Rightarrow Q$	$P \Leftrightarrow Q$
<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>true</i>
<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>false</i>
<i>true</i>	<i>false</i>	<i>false</i>	<i>false</i>	<i>true</i>	<i>false</i>	<i>false</i>
<i>true</i>	<i>true</i>	<i>false</i>	<i>true</i>	<i>true</i>	<i>true</i>	<i>true</i>

We have 3 possible worlds for $P \Rightarrow Q = \text{true}$

Propositional logic: Semantics

Rules for evaluating truth with respect to a model:

- $\neg P$ is true iff P is false
 - $P \wedge Q$ is true iff P is true and Q is true
 - $P \vee Q$ is true iff P is true or Q is true
 - $P \Rightarrow Q$ is true iff P is false or Q is true
-
- Sentence
- Model

Logical equivalence

Two sentences are **logically equivalent** iff (read if, and only if) they are true in same models

$$\begin{aligned}(\alpha \wedge \beta) &\equiv (\beta \wedge \alpha) && \text{commutativity of } \wedge \\(\alpha \vee \beta) &\equiv (\beta \vee \alpha) && \text{commutativity of } \vee \\((\alpha \wedge \beta) \wedge \gamma) &\equiv (\alpha \wedge (\beta \wedge \gamma)) && \text{associativity of } \wedge \\((\alpha \vee \beta) \vee \gamma) &\equiv (\alpha \vee (\beta \vee \gamma)) && \text{associativity of } \vee \\\neg(\neg\alpha) &\equiv \alpha && \text{double-negation elimination} \\(\alpha \Rightarrow \beta) &\equiv (\neg\beta \Rightarrow \neg\alpha) && \text{contraposition} \\(\alpha \Rightarrow \beta) &\equiv (\neg\alpha \vee \beta) && \text{implication elimination} \\(\alpha \Leftrightarrow \beta) &\equiv ((\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha)) && \text{biconditional elimination} \\\neg(\alpha \wedge \beta) &\equiv (\neg\alpha \vee \neg\beta) && \text{de Morgan} \\\neg(\alpha \vee \beta) &\equiv (\neg\alpha \wedge \neg\beta) && \text{de Morgan} \\(\alpha \wedge (\beta \vee \gamma)) &\equiv ((\alpha \wedge \beta) \vee (\alpha \wedge \gamma)) && \text{distributivity of } \wedge \text{ over } \vee \\(\alpha \vee (\beta \wedge \gamma)) &\equiv ((\alpha \vee \beta) \wedge (\alpha \vee \gamma)) && \text{distributivity of } \vee \text{ over } \wedge\end{aligned}$$

Entailment

- **Entailment** means that a sentence **follows from** the premises contained in the knowledge base:

$$KB \models \alpha$$

- The knowledge base KB entails sentence α *iff* α is true in all models where KB is true
 - E.g., KB with $x = 0$ entails sentence $x * y = 0$
- Tests for entailment
 - $KB \models \alpha$ iff $(KB \Rightarrow \alpha)$ is *valid*
 - $KB \models \alpha$ iff $(KB \wedge \neg \alpha)$ is *unsatisfiable*

Inference

- **Logical inference:** a procedure for generating sentences that follow from a knowledge base KB
- An inference procedure is **sound** if it derives a sentence α iff $KB \models \alpha$. I.e, it only derives **true sentences**.
- An inference procedure is **complete** if it can derive **all** α for which $KB \models \alpha$.

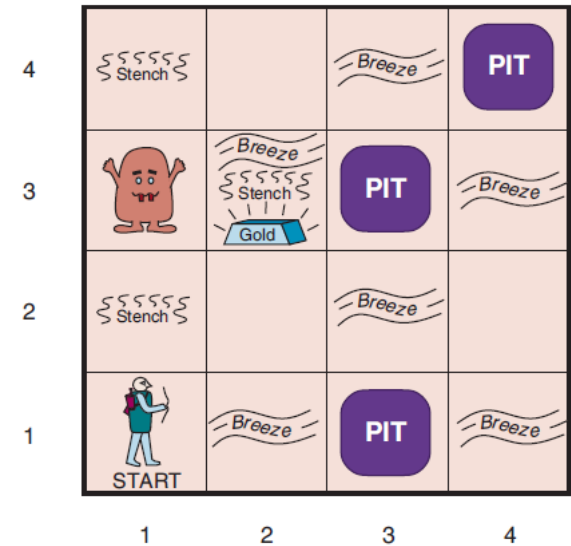
Inference

- How can we check whether a sentence α is entailed by KB?
- How about we **enumerate all possible models of the KB** (truth assignments of all its symbols), and check that α is true in every model in which KB is true?
 - Is this sound (if an answer is produced, it is correct)?
 - Is this complete (guaranteed to produce the correct answer)?
 - Problem: if KB contains n symbols, the truth table will be of size 2^n
- Better idea: use ***inference rules***, or sound procedures to generate new sentences or *conclusions* given the *premises* in the KB

Complexity of inference

- Propositional inference is ***co-NP-complete***
 - *Complement* of the SAT problem: $\alpha \models \beta$ if and only if the sentence $\alpha \wedge \neg \beta$ is *unsatisfiable*
 - Every known inference algorithm has worst-case exponential running time
- Efficient inference possible for restricted cases
 - e.g., Horn clauses are disjunctions of literals with at most one positive literal.

Example: Wumpus World



1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK			
1,1	2,1	3,1	4,1
A			
OK	OK		

(a)

A = Agent
B = Breeze
G = Glitter, Gold
OK = Safe square
P = Pit
S = Stench
V = Visited
W = Wumpus

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
OK	P?		
1,1	2,1	3,1	4,1
V	A	P?	
OK	B		
	OK		

(b)

Example: Wumpus World

Initial KB needs to contain rules like these for each square:

$$Breeze(1,1) \Leftrightarrow Pit(1,2) \vee Pit(2,1)$$

$$Breeze(1,2) \Leftrightarrow Pit(1,1) \vee Pit(1,3) \vee Pit(2,2)$$

$$Stench(1,1) \Leftrightarrow W(1,2) \vee W(2,1)$$

...

Percepts at (1,1) are no breeze or stench. Add the following facts to the KB:

$$\neg Breeze(1,1)$$

$$\neg Stench(1,1)$$

Inference will tell us that the following facts are entailed:

$$\neg Pit(1,2), \neg Pit(2,1), \neg W(1,2), \neg W(2,1)$$

This means that (1,2) and (2,1) are safe.

Initial KB needs to contain rules like these for each square:

$$Breeze(1,1) \Leftrightarrow Pit(1,2) \vee Pit(2,1)$$

$$Breeze(1,2) \Leftrightarrow Pit(1,1) \vee Pit(1,3) \vee Pit(2,2)$$

$$Stench(1,1) \Leftrightarrow W(1,2) \vee W(2,1)$$

...

Percepts at (1,1) are no breeze or stench. Add the following facts to the KB:

$$\neg Breeze(1,1)$$

$$\neg Stench(1,1)$$

Inference will tell us that the following facts are entailed:

$$\neg Pit(1,2), \neg Pit(2,1), \neg W(1,2), \neg W(2,1)$$

This means that (1,2) and (2,1) are safe.

Summary

- Logical agents apply **inference** to a **knowledge base** to derive new information and make decisions
- Basic concepts of logic:
 - **syntax**: formal structure of sentences
 - **semantics**: truth of sentences in models
 - **entailment**: necessary truth of one sentence given another
 - **inference**: deriving sentences from other sentences
 - **soundness**: derivations produce only entailed sentences
 - **completeness**: derivations can produce all entailed sentences
- Resolution is complete for propositional logic
- Algorithms use forward, backward chaining, are linear in time, and complete for special clauses (definite clauses).



First-Order Logic

Limitations of propositional logic

- Suppose you want to say “All humans are mortal”
 - In propositional logic, you would need ~6.7 billion statements of the form:
MichaelIsHuman and MichaelIsMortal,
SarahIsHuman and SarahIsMortal, ...
- Suppose you want to say “Some people can run a marathon”
 - You would need a disjunction of ~6.7 billion statements:

MichaelCanRunAMarathon or ... or SarahCanRunAMarathon

Other Languages to represent Knowledge

Language	Ontological Commitment (What exists in the world)	Epistemological Commitment (What an agent believes about facts)
Propositional logic	facts	true/false/unknown
First-order logic	facts, objects, relations	true/false/unknown
Temporal logic	facts, objects, relations, times	true/false/unknown
Probability theory	facts	degree of belief $\in [0, 1]$
Fuzzy logic	facts with degree of truth $\in [0, 1]$	known interval value

- First-order Logic adds **objects** and **relations**.

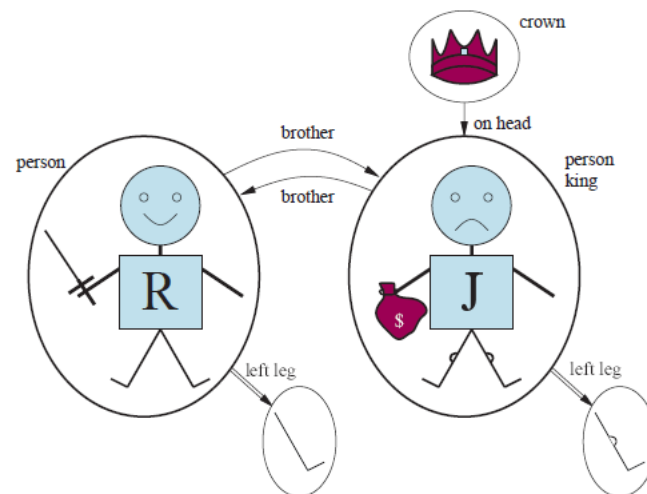
Syntax of FOL

$Sentence \rightarrow AtomicSentence \mid ComplexSentence$
 $AtomicSentence \rightarrow Predicate \mid Predicate(Term, \dots) \mid Term = Term$
 $ComplexSentence \rightarrow (Sentence)$
 $\quad \mid \neg Sentence$
 $\quad \mid Sentence \wedge Sentence$
 $\quad \mid Sentence \vee Sentence$
 $\quad \mid Sentence \Rightarrow Sentence$
 $\quad \mid Sentence \Leftrightarrow Sentence$
 $\quad \mid Quantifier Variable, \dots Sentence$

$Term \rightarrow Function(Term, \dots)$
 $\quad \mid Constant$
 $\quad \mid Variable$

$Quantifier \rightarrow \forall \mid \exists$
 $Constant \rightarrow A \mid X_1 \mid John \mid \dots$
 $Variable \rightarrow a \mid x \mid s \mid \dots$
 $Predicate \rightarrow True \mid False \mid After \mid Loves \mid Raining \mid \dots$
 $Function \rightarrow Mother \mid LeftLeg \mid \dots$

OPERATOR PRECEDENCE : $\neg, =, \wedge, \vee, \Rightarrow, \Leftrightarrow$



Objects

Relations. Predicate
is/returns True or False

Function returns an object

Universal quantification

- $\forall x P(x)$
- Example: “Everyone at SMU is smart”
 $\forall x \text{ At}(x, \text{SMU}) \Rightarrow \text{Smart}(x)$
Why not $\forall x \text{ At}(x, \text{SMU}) \wedge \text{Smart}(x)$?
- Roughly speaking, equivalent to the **conjunction** of all possible instantiations of the variable:
[$\text{At}(\text{John}, \text{SMU}) \Rightarrow \text{Smart}(\text{John})$] \wedge ...
[$\text{At}(\text{Richard}, \text{SMU}) \Rightarrow \text{Smart}(\text{Richard})$] \wedge ...
- $\forall x P(x)$ is true in a model m iff $P(x)$ is true with x being each possible object in the model

Existential quantification


- $\exists \mathbf{x} \mathbf{P}(\mathbf{x})$
- Example: “Someone at SMU is smart”
 $\exists \mathbf{x} \text{At}(\mathbf{x}, \text{SMU}) \wedge \text{Smart}(\mathbf{x})$
Why not $\exists \mathbf{x} \text{At}(\mathbf{x}, \text{SMU}) \Rightarrow \text{Smart}(\mathbf{x})$?
- Roughly speaking, equivalent to the **disjunction** of all possible instantiations:
 $[\text{At}(\text{John}, \text{SMU}) \wedge \text{Smart}(\text{John})] \vee$
 $[\text{At}(\text{Richard}, \text{SMU}) \wedge \text{Smart}(\text{Richard})] \vee \dots$
- $\exists \mathbf{x} \mathbf{P}(\mathbf{x})$ is true in a model m iff $\mathbf{P}(\mathbf{x})$ is true with \mathbf{x} being some possible object in the model

Inference in FOL

- Inference is complicated.
1. **Reduction to propositional logic** and then use propositional logic inference.
 2. **Directly do inference on FOL (or a subset like definite clauses)**
 - Unification: Combine two sentences into one.
 - Forward Chaining for FOL
 - Backward Chaining for FOL
 - Logical programming (e.g., Prolog)

Limitations of Logic

- What if we cannot collect enough knowledge to make a decision
e.g., the environment if only partially observable
- What about facts about the future?
grade(Michael, AI, this_semester)



Language	Ontological Commitment (What exists in the world)	Epistemological Commitment (What an agent believes about facts)
Propositional logic	facts	true/false/unknown
First-order logic	facts, objects, relations	true/false/unknown
Temporal logic	facts, objects, relations, times	true/false/unknown
Probability theory	facts	degree of belief $\in [0, 1]$
Fuzzy logic	facts with degree of truth $\in [0, 1]$	known interval value